


Phase 4 : Local Kubernetes Deployment - Guide & Installation

Phase 4 = Taking your working app (Phase 3) and packing it into boxes (Docker), placing those boxes into an automated factory (Kubernetes), and controlling the factory with a remote control (Helm).

NOTE: Install everything first and then create in detail specifications to proceed further! This is just a guide to help you understand what is happening in Phase 4!

Mental Model (IMPORTANT – memorize this)


Think in **4 layers**:



We'll now walk **top → bottom**, step by step.

STEP 0: What you already have (Phase 3)

You already built this 🤝



Right now:

- It runs **locally**
- Manually started
- Not scalable
- Not production-safe

Phase 4's job is to **professionalize** this.

STEP 1: Docker — “Put everything in boxes”

Why Docker?

Because Kubernetes **ONLY understands containers**.

So first:

👉 We turn frontend and backend into Docker images

Backend Docker flow (FastAPI)

FastAPI Code

```
↓  
Dockerfile
```

↓

Docker Image (todo-backend)

Inside that image:

- Python 3.12
- Your FastAPI app
- All dependencies
- Runs as non-root user
- Starts using: fastapi run

Result:

“This backend can run anywhere. Same behavior. Always.”

Frontend Docker flow (Next.js)

Next.js Code



Dockerfile



Docker Image (todo-frontend)

Important idea:

- Build happens once
- Output is **standalone**
- Final image is small & fast

Result:

“This frontend is lightweight and production-ready.”

🔑 Key idea of Step 1

After Docker:

- Your app is no longer “code”
- It is now **two images**:
 - todo-backend
 - todo-frontend

INSTALLATION:

```
**Docker** (24+)
```
Verify installation
docker --version
```
Install: https://docs.docker.com/get-docker/
```

STEP 2: Kubernetes — “Run the boxes automatically”


Docker gives boxes.

Kubernetes decides:

- how many boxes
- when to restart
- who gets traffic
- what happens if one dies

Kubernetes flow (VERY IMPORTANT)

Docker Images



Let's break that slowly.

2.1 Pods — “Run the container”

A **Pod** = one running container.

But:

- ✗ Pods die
- ✗ Pods restart
- ✗ Pods are not stable

So we **never use Pods directly**.

2.2 Deployments — “Keep pods alive”

Deployment says:

- “I want 2 backends”
- “If one dies, replace it”
- “Update safely”

So you have:

```
Backend Deployment  
└── Pod 1 (FastAPI)  
└── Pod 2 (FastAPI)
```

```
Frontend Deployment  
└── Pod 1 (Next.js)
```

└─ Pod 2 (Next.js)

2.3 Health checks — “Is it alive? Is it ready?”

Each backend pod exposes:

`/api/v1/health`

Kubernetes asks:

- **?** Alive? (liveness probe)
- **?** Ready for traffic? (readiness probe)

If not:

→ Kubernetes restarts or blocks traffic

This is **self-healing**.

2.4 Services — “Stable networking”

Pods change IPs.

Services do NOT.

So:

Frontend Pods



Frontend Service (stable)

Backend Pods




Backend Service (stable)

Now:

- Frontend can always talk to backend
- Even if pods restart

2.5 Ingress — “Public door to your app”

Ingress is the **front gate**.



Result:

- One domain
- Clean routing
- HTTPS support

INSTALLATION:

```
**kubectl** (1.28+)
```
Verify installation
kubectl version --client
```
Install: https://kubernetes.io/docs/tasks/tools/
```

🔒 STEP 3: Secrets & Config — “No secrets in code”

Rule:

- ✗ Never hardcode secrets
 - ✗ Never put secrets in Docker images
-

What goes where?

Secrets (private

Stored in **Kubernetes Secrets**:

- DATABASE_URL
- OPENAI_API_KEY
- BETTER_AUTH_SECRET etc... (according to your project)

Injected at **runtime**.

ConfigMaps (public-ish

Stored in **ConfigMaps**:

- API URLs
- Environment name
- Feature flags

Same image, different environment.

STEP 4: Helm — “One command controls everything”

Helm is NOT magic.

It's just:

“Templates + values”

Helm flow

Helm Chart

 |—— Backend templates

```
|-- Frontend templates  
|-- Ingress template  
└-- Values files
```

You run:

```
helm install todo-app ./helm/todo-app
```

Helm:

1. Reads values
 2. Fills templates
 3. Creates Kubernetes resources
-

Multi-environment magic ✨

Same chart:

```
values-dev.yaml      → 1 replica, cheap  
values-staging.yaml → 2 replicas  
values-prod.yaml    → more power
```

No code change.

No YAML duplication.

INSTALLATION:

```
**Helm** (3.x)  
```bash  
Verify installation
helm version
``

Install: https://helm.sh/docs/intro/install/
```

# STEP 5: Minikube — “Practice before production”

Minikube = **local Kubernetes**

---

## Local dev flow

```
Your Laptop
↓
Docker build
↓
minikube image load
↓
helm install
↓
App running locally (K8s-style)
```

Two DB options:

- Remote Neon (realistic)
- Local Postgres (offline)

## INSTALLATION:

```
Minikube (1.32+)
```bash  
# Verify installation  
minikube version  
``  
Install: https://minikube.sigs.k8s.io/docs/start/
```

STEP 6: What happens when things go wrong?

This is where Kubernetes shines.

Backend crashes?

→ Restarted automatically

Too much traffic?

→ HPA adds more backend pods

Pod killed?

→ Another replaces it

Deployment fails?

→ Helm rollback

User sessions?

→ Stored in DB, not memory



FINAL FLOWCHART (end-to-end)

Developer writes code



Docker builds images



Helm deploys to Kubernetes



Deployments create Pods



Services connect Pods



Ingress exposes app



Users access app



Kubernetes auto-heals + scales

✓ Why this Phase 4 is GOOD (in plain words)

Because after this:


- You can scale
- You can restart safely
- You can deploy confidently
- You can onboard other devs
- You can go to production

This is **how real companies run apps.**

Below is a **clean, start-to-end ASCII diagram** showing **EVERYTHING** in Phase 4 and **how it flows**.

Read it **top → bottom** like a pipeline.

✖ PHASE 4 — FULL CLOUD-NATIVE ASCII FLOW DIAGRAM



DOCKER
(Pack code + runtime into images)

todo-frontend Docker Image

- node:20-alpine
- Next.js standalone build
- non-root user

todo-backend Docker Image

- python:3.12-slim
- fastapi run
- non-root appuser

IMAGE REGISTRY / MINIKUBE




- Docker Hub / GHCR (prod)
- minikube image load (local)


HELM
(One command controls everything)

```
helm install todo-app ./helm/todo-app
```

Reads:

- values-dev.yaml / prod.yaml
- templates/*.yaml







SUPER IMPORTANT CONNECTIONS (read once)

Frontend Pod —calls—> Backend Service
Backend Pod —calls—> Neon PostgreSQL
Backend Pod —calls—> OpenAI API
Ingress —routes—> Frontend / Backend
Helm —creates—> EVERYTHING above



WHAT HAPPENS DURING FAILURE (quick mini-flow)

Backend Pod crashes

↓

 Liveness probe fails

↓

 Kubernetes kills pod

↓

 Deployment creates new pod

↓

 Service routes traffic to healthy pod

↓

 User never notices



ONE-LINE MEMORY HOOK

Docker packs it, Helm installs it, Kubernetes runs it, Ingress exposes it.